

## EXPERIMENT 9

### Aim:

To build a Logistic Regression model using the Social Network Ads dataset to predict whether a user will purchase a product based on their age and estimated salary.

### Algorithm:

1. Import necessary libraries such as NumPy, Pandas, and Scikit-learn.
2. Load the dataset using Pandas.
3. Display the dataset and inspect the first few records.
4. Extract the feature columns (Age, Estimated Salary) and label (Purchased).
5. Split the dataset into training and testing sets using `train_test_split()`.
6. Train the Logistic Regression model using the training data.
7. Evaluate the model accuracy using `score()` on both training and testing datasets.
8. Identify the best random state by iterating over multiple random states to maximize test accuracy.
9. Display the classification report to evaluate performance using precision, recall, and F1-score.

### Code:

```
import numpy as np
import pandas as pd
df = pd.read_csv('Social_Network_Ads - Social_Network_Ads.csv')
df
```

### Output:

```
User ID  Gender  Age  EstimatedSalary  Purchased
0  15624510  Male   19     19000         0
1  15810944  Male   35     20000         0
2  15668575  Female  26     43000         0
3  15603246  Female  27     57000         0
4  15804002  Male   19     76000         0
...
395 15691863  Female  46     41000         1
```

```
396 15706071 Male 51 230000 1
397 15654296 Female 50 20000 0
398 15755018 Male 36 33000 0
399 15594041 Female 49 36000 1
```

400 rows × 5 columns

---

df.head()

**Output:**

```
User ID Gender Age EstimatedSalary Purchased
0 15624510 Male 19 19000 0
1 15810944 Male 35 20000 0
2 15668575 Female 26 43000 0
3 15603246 Female 27 57000 0
4 15804002 Male 19 76000 0
```

---

```
features = df.iloc[:, [2, 3]].values
```

```
label = df.iloc[:, 4].values
```

```
features
```

**Output:**

```
array([[ 19, 19000],
       [ 35, 20000],
       [ 26, 43000],
       [ 27, 57000],
       [ 19, 76000],
       [ 27, 58000],
       [ 27, 84000],
       [ 32, 150000],
       [ 25, 33000],
       [ 35, 65000],
       [ 26, 80000],
       [ 26, 52000],
```

```
[ 20, 86000],  
[ 32, 18000],  
[ 18, 82000],  
[ 29, 80000],  
[ 47, 25000],  
[ 45, 26000],  
[ 46, 28000],  
[ 48, 29000],  
...  
[ 49, 36000]])
```

---

label

**Output:**

```
array([[0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,  
       1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
       0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
       0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
       0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,  
       0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,  
       0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
       0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,  
       0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0,  
       1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,  
       1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1,  
       0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1,  
       1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,  
       0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,  
       1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1,  
       0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1,  
       1, 1, 0, 1]])
```

---

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

for i in range(1, 401):
    x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=i)
    model = LogisticRegression()
    model.fit(x_train, y_train)
    train_score = model.score(x_train, y_train)
    test_score = model.score(x_test, y_test)
    if test_score > train_score:
        print("Test {} Train{} Random State {}".format(test_score, train_score, i))
```

**Output:**

```
Test 0.9 Train0.840625 Random State 4
Test 0.8625 Train0.85 Random State 5
Test 0.8625 Train0.859375 Random State 6
Test 0.8875 Train0.8375 Random State 7
Test 0.8625 Train0.8375 Random State 9
...
Test 0.95 Train0.81875 Random State 352
Test 0.9625 Train0.81875 Random State 217
Test 0.9375 Train0.821875 Random State 364
```

---

```
x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=42)
finalModel = LogisticRegression()
finalModel.fit(x_train, y_train)
```

---

```
print(finalModel.score(x_train, y_train))
print(finalModel.score(x_test, y_test))
```

**Output:**

```
0.8375
```

0.8875

---

```
from sklearn.metrics import classification_report  
print(classification_report(label, finalModel.predict(features)))
```

**Output:**

	precision	recall	f1-score	support
0	0.85	0.93	0.89	257
1	0.85	0.70	0.77	143
accuracy			0.85	400
macro avg	0.85	0.81	0.83	400
weighted avg	0.85	0.85	0.84	400

**Result:**

Thus, the Logistic Regression model was successfully trained and tested using the Social Network Ads dataset.

The model achieved an accuracy of **88.75%** on the test data and **85% overall**, effectively predicting whether users are likely to purchase a product based on their age and estimated salary.